RAJALAKSHMI ENGINEERING COLLEGE

(An Autonomous Institution)

RAJALAKSHMI NAGAR, THANDALAM- 602 105

CS19P18 - DEEP LEARNING CONCEPTS

LABORATORY RECORD NOTEBOOK

NAME: SURANJANA S

YEAR/SEMESTER: IV / VII

BRANCH: COMPUTER SCIENCE AND ENGINEERING

REGISTER NO: 220701293

COLLEGE ROLL NO: 2116220701293

ACADEMIC YEAR: 2025 -2026

RAJALAKSHMI ENGINEERING COLLEGE

(An Autonomous Institution)

RAJALAKSHMI NAGAR, THANDALAM- 602 105

BONAFIDE CERTIFICATE

NAME: SURANJANA .S **BRANCH/SECTION:** COMPUTER SCIENCE AND ENGINEERING/D **ACADEMIC YEAR:** 2025 -2026   **SEMESTER:** SEVEN

REGISTER NO: | **220701293** |

Certified that this is a Bonafide record of work done by the above student in the **CS19P18 -   DEEP LEARNING CONCEPTS**

during the year 2025 - 2026

Signature of Faculty In-charge

Submitted for the Practical Examination Held on: …………………………………………...

Internal Examiner   **External Examiner**

INDEX

INSTALLATION AND CONFIGURATION OF TENSORFLOW

Aim:

To install and configure TensorFlow in the anaconda environment in Windows 10.

Procedure:

1. Download Anaconda Navigator and install.

2. Open Anaconda prompt

3. Create a new environment dlc with python 3.7 using the following
   command: conda create -n dlc python=3.7

4. Activate newly created environment dlc using the following
   command: conda activate dlc

5. In dlc prompt, install tensorflow using the following
   command: pip install tensorflow

6. Next install Tensorflow-datasets using the following
   command: pip install tensorflow-datasets

7. Install scikit-learn package using the following
   command: pip install scikit-learn

8. Install pandas package using the following
   command: pip install pandas

9. Lastly, install jupyter notebook
   pip install jupyter
   notebook

10. Open jupyter notebook by typing the following in dlc
    prompt: jupyter notebook

11. Click create new and then choose python 3 (ipykernel)

12. Give the name to the file

13. Type the code and click Run button to execute (eg. Type import tensorflow and then run)

EX NO: 1   CREATE A NEURAL NETWORK TO RECOGNIZE HANDWRITTEN

DATE:14/07/2025   DIGITS USING MNIST DATASET

Aim:

To build a handwritten digit's recognition with MNIST dataset.

Procedure:

1. Download and load the MNIST dataset.

2. Perform analysis and preprocessing of the dataset.

3. Build a simple neural network model using Keras/TensorFlow.

4. Compile and fit the model.

5. Perform prediction with the test dataset.

6. Calculate performance metrics.

Code:

```
import numpy as np
import tensorflow as
tf
from tensorflow import keras

from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score
# Generate a synthetic dataset

X, y = make_classification(n_samples=1000, n_features=20,
random_state=42) # Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42) # Standardize features (optional but often beneficial)
scaler = StandardScaler()

X_train = scaler.fit_transform(X_train)
```

```python
X_test = scaler.transform(X_test
# Define the model
model = keras.Sequential([
    keras.layers.Input(shape=(X_train.shape[1],)), # Input layer
    keras.layers.Dense(64, activation='relu'),    # Hidden layer with 64 neurons and ReLU
    activation keras.layers.Dense(1, activation='sigmoid') # Output layer with 1 neuron and
    sigmoid activation
)

# Train the model

history = model.fit(X_train, y_train, epochs=10, batch_size=32, validation_split=0.1)
# Evaluate the model on the test set
y_pred = model.predict(X_test)
y_pred_classes = (y_pred >
0.5).astype(int) # Calculate accuracy on
the test set
accuracy = accuracy_score(y_test,
y_pred_classes) # Calculate test loss
test_loss = model.evaluate(X_test, y_test)
print(f"Test accuracy: {accuracy * 100:.2f}%")
print(f"Test loss: {test_loss[0]:.4f}")
```

Output:

```
Epoch 1/10
192/192 [==============================] - 5s 17ms/step - loss: 0.3739 - accuracy: 0.8950 - val_loss: 0.1801 - val_accuracy: 0.
9480
Epoch 2/10
192/192 [==============================] - 3s 14ms/step - loss: 0.1492 - accuracy: 0.9562 - val_loss: 0.1261 - val_accuracy: 0.
9635
Epoch 3/10
192/192 [==============================] - 2s 13ms/step - loss: 0.0980 - accuracy: 0.9714 - val_loss: 0.1129 - val_accuracy: 0.
9676
Epoch 4/10
192/192 [==============================] - 2s 11ms/step - loss: 0.0711 - accuracy: 0.9795 - val_loss: 0.0962 - val_accuracy: 0.
9709
Epoch 5/10
192/192 [==============================] - 2s 10ms/step - loss: 0.0543 - accuracy: 0.9844 - val_loss: 0.0914 - val_accuracy: 0.
9725
Epoch 6/10
192/192 [==============================] - 2s 11ms/step - loss: 0.0402 - accuracy: 0.9888 - val_loss: 0.0866 - val_accuracy: 0.
9737
Epoch 7/10
192/192 [==============================] - 2s 12ms/step - loss: 0.0301 - accuracy: 0.9920 - val_loss: 0.0871 - val_accuracy: 0.
9750
Epoch 8/10
192/192 [==============================] - 2s 12ms/step - loss: 0.0245 - accuracy: 0.9931 - val_loss: 0.0840 - val_accuracy: 0.
9762
Epoch 9/10
192/192 [==============================] - 2s 12ms/step - loss: 0.0180 - accuracy: 0.9956 - val_loss: 0.0878 - val_accuracy: 0.
9760
Epoch 10/10
192/192 [==============================] - 2s 11ms/step - loss: 0.0149 - accuracy: 0.9963 - val_loss: 0.0858 - val_accuracy: 0.
9755
```

```python
In [1]: import numpy as np
        import matplotlib.pyplot as plt
        import tensorflow as tf
        from tensorflow.keras.datasets import mnist
        from tensorflow.keras.models import Sequential
        from tensorflow.keras.layers import Dense
        from tensorflow.keras.utils import to_categorical
```

```python
In [2]: feature_vector_length =784
        num_classes=10
```

```python
In [3]: (X_train, Y_train), (X_test, Y_test) = mnist.load_data()
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11490434/11490434 [==============================] - 2s 0us/step
```

```python
In [4]: input_shape = (feature_vector_length)
        print(f'Feature shape: {input_shape}')
```

```
Feature shape: 784
```

```python
In [5]: X_train = X_train.reshape(X_train.shape[0], feature_vector_length)
        X_test = X_test.reshape(X_test.shape[0], feature_vector_length)
```

```python
In [6]: X_train = X_train.astype('float32') / 255
        X_test = X_test.astype('float32') / 255
        Y_train = to_categorical(Y_train, num_classes)
        Y_test = to_categorical(Y_test, num_classes)
```


Sample 1 - Predicted: 7,Actual: 7   Sample 2 - Predicted: 2,Actual: 2   Sample 3 - Predicted: 1,Actual: 1

Sample 4 - Predicted: 0,Actual: 0   Sample 5 - Predicted: 4,Actual: 4
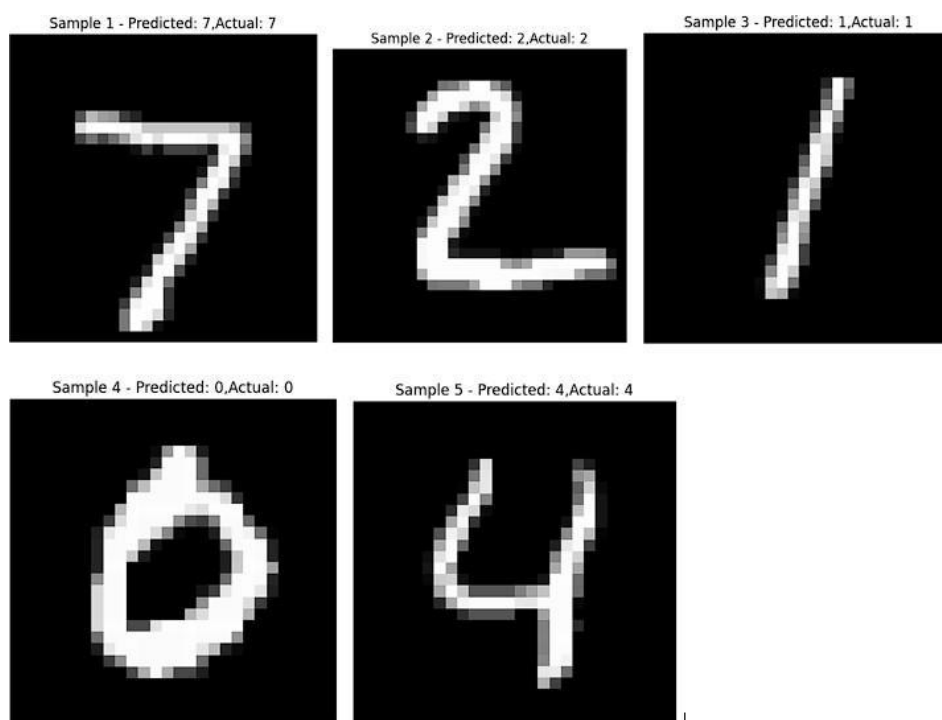
Result:

Thus, the implementation to build a simple neural network using Keras/TensorFlow has been successfully executed.

EX NO:2   **BUILD A CONVOLUTIONAL NEURAL NETWORK**

DATE:21/07/2025   **USING KERAS/TENSORFLOW**

Aim:

To implement a Convolutional Neural Network (CNN) using Keras/TensorFlow to recognize and classify handwritten digits from the MNIST dataset with high accuracy.

Procedure:

1. Import required libraries (TensorFlow/Keras, NumPy, etc.).

2. Load the MNIST dataset from Keras.

3. Normalize and reshape the image data.

4. Convert labels to one-hot encoded vectors.

5. Build a CNN model with Conv2D, MaxPooling, Flatten, and Dense layers.

6. Compile the model using categorical crossentropy and Adam optimizer.

7. Train the model on training data.

8. Evaluate the model on test data.

9. Display accuracy and predictions.

Code:

```
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from tensorflow.keras.datasets import mnist
import matplotlib.pyplot as plt
import numpy as np

(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
train_images = train_images / 255.0
test_images = test_images / 255.0
train_images = train_images.reshape(-1, 28, 28, 1)
test_images = test_images.reshape(-1, 28, 28, 1)
model = Sequential([
Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
MaxPooling2D((2, 2)),
Conv2D(64, (3, 3), activation='relu'),
MaxPooling2D((2, 2)),
Flatten(),
Dense(64, activation='relu'),
Dropout(0.5),
Dense(10, activation='softmax')
])
```

```python
model.compile(optimizer='adam',
loss='sparse_categorical_crossentropy',
metrics=['accuracy'])

history = model.fit(train_images, train_labels,
epochs=5,
batch_size=64,
validation_split=0.2)

test_loss, test_acc = model.evaluate(test_images, test_labels)
print(f"\n Test accuracy: {test_acc:.4f}")
print(f" Test loss: {test_loss:.4f}")

plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Train Accuracy', marker='o')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy', marker='o')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.grid(True)

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Train Loss', marker='o')
plt.plot(history.history['val_loss'], label='Validation Loss', marker='o')
plt.title('Training and Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()

predictions = model.predict(test_images)
predicted_labels = np.argmax(predictions, axis=1)

num_samples = 10
plt.figure(figsize=(15, 4))

for i in range(num_samples):
plt.subplot(1, num_samples, i + 1)
plt.imshow(test_images[i].reshape(28, 28), cmap='gray')
plt.title(f"Pred: {predicted_labels[i]}\nTrue: {test_labels[i]}")
```
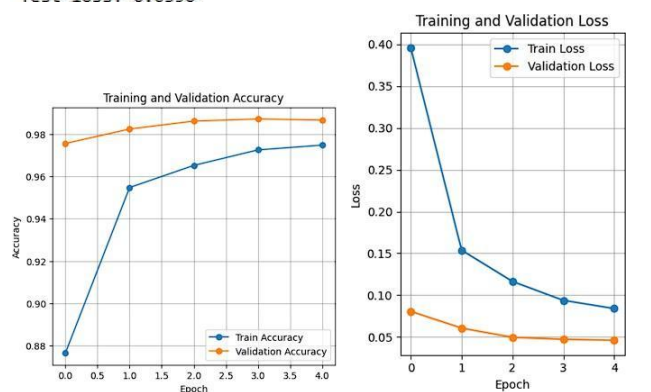
plt.axis('off')
plt.suptitle("Sample Predictions on Test Images", fontsize=16)
plt.show()

Output:

```
Epoch 1/5
750/750 [==============================] - 30s 39ms/step - loss: 0.3961 - accuracy: 0.8765 - val_loss: 0.0806 - val_accuracy:
0.9756
Epoch 2/5
750/750 [==============================] - 26s 35ms/step - loss: 0.1538 - accuracy: 0.9548 - val_loss: 0.0606 - val_accuracy:
0.9824
Epoch 3/5
750/750 [==============================] - 30s 39ms/step - loss: 0.1163 - accuracy: 0.9652 - val_loss: 0.0495 - val_accuracy:
0.9862
Epoch 4/5
750/750 [==============================] - 27s 36ms/step - loss: 0.0937 - accuracy: 0.9725 - val_loss: 0.0472 - val_accuracy:
0.9872
Epoch 5/5
750/750 [==============================] - 26s 35ms/step - loss: 0.0840 - accuracy: 0.9748 - val_loss: 0.0460 - val_accuracy:
0.9867
313/313 [==============================] - 2s 5ms/step - loss: 0.0390 - accuracy: 0.9880
```

```
Test accuracy: 0.9880
Test loss: 0.0390
```


Result:

Thus, the Convolution Neural Network (CNN) using Keras / Tensorflow to recognize and classify handwritten digits from MNIST dataset has been implemented successfully.

EX NO: 3   IMAGE CLASSIFICATION ON CIFAR-10 DATASET USING CNN

DATE:28/07/2025

Aim:

To build a Convolutional Neural Network (CNN) model for classifying images from the CIFAR-10 dataset into one of the ten categories such as airplanes, cars, birds, cats, etc.

Procedure:

1. Download and load the CIFAR-10 dataset using Keras/TensorFlow.

2. Visualize and analyze sample images from the

dataset. 3, Preprocess the data:

- Normalize the pixel values (divide by 255)

- Convert class labels to one-hot encoded format

4. Build a CNN model using Keras/TensorFlow:

- Include convolutional, pooling, flatten, and dense layers.

5. Compile the model with a suitable loss function and optimizer.

6. Train the model using training data and validate using test data.

7. Evaluate the model using accuracy and loss on the test dataset.

8. Perform predictions on new/unseen CIFAR-10 images.

9 Visualize prediction results with sample images and predicted labels.

Code:

```
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.cifar10.load_data()
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0
y_train = tf.keras.utils.to_categorical(y_train, 10)
y_test = tf.keras.utils.to_categorical(y_test, 10)
model = tf.keras.Sequential()
model.add(tf.keras.layers.Conv2D(32, (3,3), activation='relu',
input_shape=(32,32,3))) model.add(tf.keras.layers.MaxPooling2D((2,2)))
model.add(tf.keras.layers.Conv2D(64, (3,3), activation='relu'))
model.add(tf.keras.layers.MaxPooling2D((2,2)))
model.add(tf.keras.layers.Conv2D(64, (3,3), activation='relu'))
model.add(tf.keras.layers.Flatten())
model.add(tf.keras.layers.Dense(64, activation='relu'))
model.add(tf.keras.layers.Dense(10, activation='softmax'))
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
model.fit(x_train, y_train, epochs=10, batch_size=64, validation_split=0.2)
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
```

'dog', 'frog', 'horse', 'ship', 'truck']

```
index = int(input("Enter an index (0 to 9999) for test image: "))
if index < 0 or index >= len(x_test):
print("Invalid index. Using index 0 by default.")
index = 0
test_image = x_test[index]
true_label = np.argmax(y_test[index])
prediction = model.predict(np.expand_dims(test_image, axis=0))
predicted_label = np.argmax(prediction)
plt.figure(figsize=(4, 4))
resized_image = tf.image.resize(test_image, [128, 128])
plt.imshow(resized_image)
plt.axis('off')
plt.title(f"Predicted: {class_names[predicted_label]}\nActual: {class_names[true_label]}")
plt.show()
```

Output:

```
Epoch 1/10
625/625 [==============================] - 58s 87ms/step - loss: 1.6801 - accuracy: 0.3846 - val_loss: 1.4341 - val_accuracy:
0.4803
Epoch 2/10
625/625 [==============================] - 37s 60ms/step - loss: 1.3153 - accuracy: 0.5284 - val_loss: 1.3005 - val_accuracy:
0.5388
Epoch 3/10
625/625 [==============================] - 36s 58ms/step - loss: 1.1663 - accuracy: 0.5846 - val_loss: 1.1370 - val_accuracy:
0.6014
Epoch 4/10
625/625 [==============================] - 38s 61ms/step - loss: 1.0629 - accuracy: 0.6249 - val_loss: 1.0984 - val_accuracy:
0.6178
Epoch 5/10
625/625 [==============================] - 41s 65ms/step - loss: 0.9991 - accuracy: 0.6480 - val_loss: 1.0476 - val_accuracy:
0.6379
Epoch 6/10
625/625 [==============================] - 38s 61ms/step - loss: 0.9348 - accuracy: 0.6720 - val_loss: 0.9795 - val_accuracy:
0.6598
Epoch 7/10
625/625 [==============================] - 38s 60ms/step - loss: 0.8764 - accuracy: 0.6970 - val_loss: 1.0013 - val_accuracy:
0.6547
Epoch 8/10
625/625 [==============================] - 38s 61ms/step - loss: 0.8338 - accuracy: 0.7096 - val_loss: 0.9313 - val_accuracy:
0.6770
Epoch 9/10
625/625 [==============================] - 39s 62ms/step - loss: 0.7943 - accuracy: 0.7242 - val_loss: 0.9243 - val_accuracy:
0.6856
Epoch 10/10
625/625 [==============================] - 37s 60ms/step - loss: 0.7588 - accuracy: 0.7362 - val_loss: 0.8994 - val_accuracy:
0.6986
```

Predicted: frog
Actual: frog

Result

Thus, the Convolution Neural Network (CNN) model for classifying images from CIFAR-10 dataset is implemented successfully.

Ex No: 4 TRANSFER LEARNING WITH CNN AND VISUALIZATION
DATE:04/08/2025

Aim:

To build a convolutional neural network with transfer learning and perform visualization

Procedure:

1. Download and load the dataset.

2. Perform analysis and preprocessing of the dataset.

3. Build a simple neural network model using Keras/TensorFlow.

4. Compile and fit the model.

5. Perform prediction with the test dataset.

6. Calculate performance metrics.

Code:

```
conda install -c conda-forge python-graphviz -y
import tensorflow as tf
from tensorflow.keras.applications import VGG16
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten, Dropout
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.utils import plot_model
import matplotlib.pyplot as plt
import numpy as np
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
x_train = x_train / 255.0
x_test = x_test / 255.0
vgg_base = VGG16(weights='imagenet', include_top=False, input_shape=(32, 32, 3))
for layer in vgg_base.layers:
layer.trainable = False
model = Sequential()
model.add(vgg_base)
model.add(Flatten())
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(10, activation='softmax'))
model.compile(optimizer=Adam(learning_rate=0.0001),
loss='sparse_categorical_crossentropy',
metrics=['accuracy'])
plot_model(model, to_file='cnn.png', show_shapes=True,
```

```python
                      show_layer_names=True, dpi=300)
plt.figure(figsize=(20, 20))
img = plt.imread('cnn.png')
plt.imshow(img)
plt.axis('off')
plt.show()
history = model.fit(x_train, y_train,
epochs=10,
batch_size=32,
validation_split=0.2)

test_loss, test_acc = model.evaluate(x_test, y_test)
print(f'Test Loss: {test_loss:.4f}')
print(f'Test Accuracy: {test_acc * 100:.2f}%')
plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Model Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Model Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.tight_layout()
plt.show()
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
'dog', 'frog', 'horse', 'ship', 'truck']
sample = x_test[0].reshape(1, 32, 32, 3)
prediction = model.predict(sample)
predicted_class = class_names[np.argmax(prediction)]

plt.imshow(x_test[0])
plt.title(f"Predicted: {predicted_class}")
plt.axis('off')
plt.show()
```

Output:

vgg16_input	input:	[(None, 32, 32, 3)]
InputLayer	output:	[(None, 32, 32, 3)]

vgg16	input:	(None, 32, 32, 3)
Functional	output:	(None, 1, 1, 512)

flatten	input:	(None, 1, 1, 512)
Flatten	output:	(None, 512)

dense	input:	(None, 512)
Dense	output:	(None, 512)

dropout	input:	(None, 512)
Dropout	output:	(None, 512)

dense_1	input:	(None, 512)
Dense	output:	(None, 10)

```
Epoch 1/10
1250/1250 [==============================] - 231s 182ms/step - loss: 1.7748 - accuracy: 0.3727 - val_loss: 1.4674 - val_accurac
y: 0.4959
Epoch 2/10
1250/1250 [==============================] - 193s 154ms/step - loss: 1.4665 - accuracy: 0.4920 - val_loss: 1.3556 - val_accurac
y: 0.5322
Epoch 3/10
1250/1250 [==============================] - 187s 150ms/step - loss: 1.3733 - accuracy: 0.5264 - val_loss: 1.2966 - val_accurac
y: 0.5512
Epoch 4/10
1250/1250 [==============================] - 189s 151ms/step - loss: 1.3197 - accuracy: 0.5386 - val_loss: 1.2610 - val_accurac
y: 0.5621
Epoch 5/10
1250/1250 [==============================] - 191s 153ms/step - loss: 1.2777 - accuracy: 0.5551 - val_loss: 1.2352 - val_accurac
y: 0.5739
Epoch 6/10
1250/1250 [==============================] - 190s 152ms/step - loss: 1.2474 - accuracy: 0.5683 - val_loss: 1.2154 - val_accurac
y: 0.5759
Epoch 7/10
1250/1250 [==============================] - 187s 150ms/step - loss: 1.2269 - accuracy: 0.5741 - val_loss: 1.1981 - val_accurac
y: 0.5830
Epoch 8/10
1250/1250 [==============================] - 183s 146ms/step - loss: 1.2039 - accuracy: 0.5834 - val_loss: 1.1839 - val_accurac
y: 0.5864
Epoch 9/10
1250/1250 [==============================] - 177s 142ms/step - loss: 1.1866 - accuracy: 0.5887 - val_loss: 1.1735 - val_accurac
y: 0.5900
Epoch 10/10
1250/1250 [==============================] - 175s 140ms/step - loss: 1.1699 - accuracy: 0.5943 - val_loss: 1.1672 - val_accurac
y: 0.5910
```


Predicted: frog
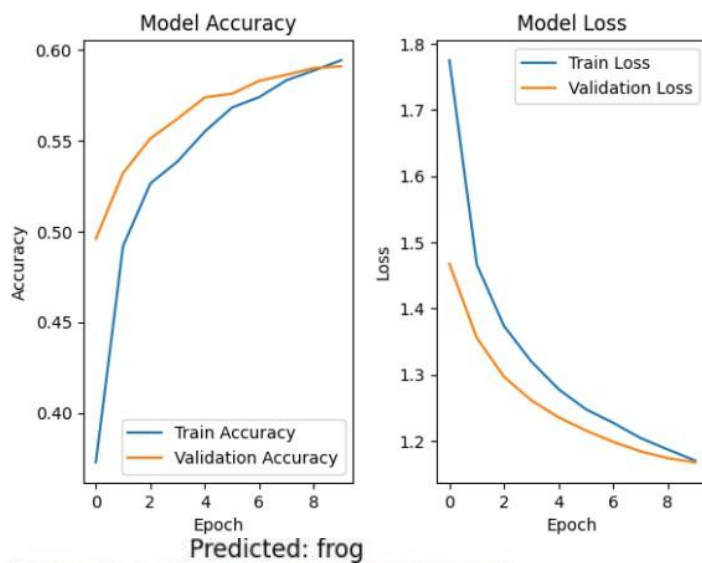
Result

Thus, the Convolution Neural Network (CNN) with transfer learning and perform visualization has been implemented successfully

EX NO: 5   BUILD A RECURRENT NEURAL NETWORK (RNN) USING

DATE:25/08/2025   **KERAS/TENSORFLOW**

Aim:

To build a recurrent neural network with Keras/TensorFlow.

Procedure:

1. Download and load the dataset.

2. Perform analysis and preprocessing of the dataset.

3. Build a simple neural network model using Keras/TensorFlow.

4. Compile and fit the model.

5. Perform prediction with the test dataset.

6. Calculate performance metrics.

Code:

```
import numpy as np
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import SimpleRNN,
Dense from sklearn.metrics import r2_score
np.random.seed(0)
seq_length = 10
num_samples = 1000
X = np.random.randn(num_samples, seq_length, 1)
y = X.sum(axis=1) + 0.1 * np.random.randn(num_samples, 1)
split_ratio = 0.8
split_index = int(split_ratio * num_samples)
X_train, X_test = X[:split_index], X[split_index:]
y_train, y_test = y[:split_index], y[split_index:]
model = Sequential()
model.add(SimpleRNN(units=50, activation='relu', input_shape=(seq_length, 1)))
model.add(Dense(units=1))
model.compile(optimizer='adam', loss='mean_squared_error')
model.summary()
batch_size = 30
epochs = 50 # Reduced epochs for quick demonstration
history = model.fit(
X_train, y_train,
batch_size=batch_size,
epochs=epochs,
validation_split=0.2
)
test_loss = model.evaluate(X_test, y_test)
print(f'Test Loss: {test_loss:.4f}')
```

```
y_pred = model.predict(X_test)
r2 = r2_score(y_test, y_pred)
print(f'Test Accuracy (R^2): {r2:.4f}')

new_data = np.random.randn(5, seq_length, 1)
predictions = model.predict(new_data)
print("Predictions for new data:")
print(predictions)
```

Output:

```
Model: "sequential"
_____
 Layer (type)                Output Shape              Param #
=================================================================
 simple_rnn (SimpleRNN)      (None, 50)                2600

 dense (Dense)               (None, 1)                 51

=================================================================
Total params: 2,651
Trainable params: 2,651
Non-trainable params: 0
_____
```

```
Epoch 1/50
22/22 [==============================] - 2s 23ms/step - loss: 8.7454
- val_loss: 6.3263
Epoch 2/50
22/22 [==============================] - 0s 4ms/step - loss: 5.8837
- val_loss: 3.7798
Epoch 3/50
22/22 [==============================] - 0s 5ms/step - loss: 3.7728
- val_loss: 2.3105
Epoch 4/50
22/22 [==============================] - 0s 5ms/step - loss: 1.7141
- val_loss: 0.5373
Epoch 5/50
22/22 [==============================] - 0s 4ms/step - loss: 0.2878
- val_loss: 0.2417
Epoch 6/50
22/22 [==============================] - 0s 4ms/step - loss: 0.1304
- val_loss: 0.1146
Epoch 7/50
```

```
1/1 [==============================] - 0s 20ms/step
Predictions for new data:
[[ 1.5437698]
 [ 0.4290885]
 [-2.1180325]
 [-0.5443404]
 [-3.8416493]]
```

Result:

Thus, the Recurrent Neural Network (RNN) has been implemented using Tensorflow.

EX NO: 6 SENTIMENT CLASSIFICATION OF TEXT USING RNN

DATE:15/09/2025

Aim:

To implement a Recurrent Neural Network (RNN) using Keras/TensorFlow for classifying the sentiment of text data (e.g., movie reviews) as positive or negative.

Procedure:

1. Import necessary libraries.

2. Load and preprocess the text dataset (e.g., IMDb).

3. Pad sequences and prepare labels.

4. Build an RNN model with Embedding and SimpleRNN layers.

5. Compile the model with loss and optimizer.

6. Train the model on training data.

7. Evaluate the model on test data.

8. Predict sentiment for new inputs

Code:

```
import numpy as np
from tensorflow.keras.datasets import imdb
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, SimpleRNN, Dense
max_words = 5000
max_len = 200
(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=max_words)
X_train = pad_sequences(x_train, maxlen=max_len)
X_test = pad_sequences(x_test, maxlen=max_len)
model = Sequential()
model.add(Embedding(input_dim=max_words, output_dim=32, input_length=max_len))
model.add(SimpleRNN(32))
model.add(Dense(1, activation='sigmoid'))
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
print("Training...")
model.fit(X_train, y_train, epochs=2, batch_size=64, validation_split=0.2)
loss, acc = model.evaluate(X_test, y_test)
print(f"\nTest Accuracy: {acc:.4f}")
word_index = imdb.get_word_index()
reverse_word_index = {v: k for (k, v) in word_index.items()}
```

```
def decode_review(review):
return " ".join([reverse_word_index.get(i - 3, "?") for i in review])
sample_review = X_test[0]
prediction = model.predict(sample_review.reshape(1, -1))[0][0]
print("\nReview text:", decode_review(x_test[0]))
print("Predicted Sentiment:", "Positive " if prediction > 0.5 else "Negative ")
```

Output:

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb.npz
17464789/17464789 ──────────────── 0s 0us/step
Training...
Epoch 1/2
/usr/local/lib/python3.12/dist-packages/keras/src/layers/core/embedding.py:97: UserWarning: Argument `input_length` is deprecated. Just remove it.
  warnings.warn(
313/313 ──────────────── 21s 59ms/step - accuracy: 0.6479 - loss: 0.6143 - val_accuracy: 0.6644 - val_loss: 0.6085
Epoch 2/2
313/313 ──────────────── 17s 53ms/step - accuracy: 0.7939 - loss: 0.4496 - val_accuracy: 0.8186 - val_loss: 0.4121
782/782 ──────────────── 10s 13ms/step - accuracy: 0.8237 - loss: 0.4115

Test Accuracy: 0.8230
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb_word_index.json
1641221/1641221 ──────────────── 0s 0us/step
```

```
def decode_review(review):
    return " ".join([reverse_word_index.get(i - 3, "?") for i in review])
sample_review = X_test[0]
prediction = model.predict(sample_review.reshape(1, -1))[0][0]
print("\nReview text:", decode_review(x_test[0]))
print("Predicted Sentiment:", "Positive " if prediction > 0.5 else "Negative ")
```

```
1/1 ──────────────── 0s 195ms/step

Review text: ? please give this one a miss br br ? ? and the rest of the cast ? terrible performances the show is flat flat flat br br i don't know how michael
Predicted Sentiment: Negative
```

Result

Thus, the Recurrent Neural Network (RNN) using Keras has been implemented for classifying sentiment of text successfully.

Ex No: 7 BUILD AUTOENCODERS WITH KERAS/TENSORFLOW

DATE:22/09/2025

Aim:

To build autoencoders with Keras/TensorFlow.

Procedure:

1. Download and load the dataset.

2. Perform analysis and preprocessing of the dataset.

3. Build a simple neural network model using Keras/TensorFlow.

4. Compile and fit the model.

5. Perform prediction with the test dataset.

6. Calculate performance metrics.

Code:

```
import numpy as np
import matplotlib.pyplot as plt
from keras.layers import Input, Dense
from keras.models import Model
from keras.datasets import mnist
(x_train, _), (x_test, _) = mnist.load_data()
x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.
x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))
input_img = Input(shape=(784,))
encoded = Dense(32, activation='relu')(input_img)
decoded = Dense(784, activation='sigmoid')(encoded)
autoencoder = Model(input_img, decoded)
autoencoder.compile(optimizer='adam', loss='binary_crossentropy')
autoencoder.fit(x_train, x_train,
epochs=50,
batch_size=256,
shuffle=True,
validation_data=(x_test, x_test))
test_loss = autoencoder.evaluate(x_test, x_test)
decoded_imgs = autoencoder.predict(x_test)
threshold = 0.5
correct_predictions = np.sum(
np.where(x_test >= threshold, 1, 0) ==
np.where(decoded_imgs >= threshold, 1, 0))
total_pixels = x_test.shape[0] * x_test.shape[1]
test_accuracy = correct_predictions / total_pixels
print("Test Loss:", test_loss)
```

```
print("Test Accuracy:", test_accuracy)
n = 10
plt.figure(figsize=(20, 4))
for i in range(n):
# Display original
ax = plt.subplot(2, n, i + 1)
plt.imshow(x_test[i].reshape(28, 28))
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)
# Display reconstruction with threshold
ax = plt.subplot(2, n, i + 1 + n)
reconstruction = decoded_imgs[i].reshape(28, 28)
plt.imshow(np.where(reconstruction >= threshold, 1.0, 0.0))
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)
plt.show()
```

Output:

```
Epoch 1/50
235/235 ──────────────── 6s 18ms/step - loss: 0.3805 - val_loss: 0.1906
Epoch 2/50
235/235 ──────────────── 5s 19ms/step - loss: 0.1808 - val_loss: 0.1547
Epoch 3/50
235/235 ──────────────── 5s 19ms/step - loss: 0.1501 - val_loss: 0.1342
Epoch 4/50
235/235 ──────────────── 3s 10ms/step - loss: 0.1321 - val_loss: 0.1221
Epoch 5/50
235/235 ──────────────── 2s 9ms/step - loss: 0.1210 - val_loss: 0.1138
Epoch 6/50
235/235 ──────────────── 3s 11ms/step - loss: 0.1134 - val_loss: 0.1081
Epoch 7/50
235/235 ──────────────── 5s 9ms/step - loss: 0.1079 - val_loss: 0.1039
Epoch 8/50
235/235 ──────────────── 2s 9ms/step - loss: 0.1042 - val_loss: 0.1006
Epoch 9/50
235/235 ──────────────── 3s 9ms/step - loss: 0.1011 - val_loss: 0.0981
Epoch 10/50
235/235 ──────────────── 3s 11ms/step - loss: 0.0989 - val_loss: 0.0963
Epoch 11/50
235/235 ──────────────── 3s 12ms/step - loss: 0.0972 - val_loss: 0.0951
Epoch 12/50
235/235 ──────────────── 3s 11ms/step - loss: 0.0964 - val_loss: 0.0943
Epoch 13/50
235/235 ──────────────── 2s 10ms/step - loss: 0.0954 - val_loss: 0.0938
Epoch 14/50
235/235 ──────────────── 2s 10ms/step - loss: 0.0950 - val_loss: 0.0934
Epoch 15/50
235/235 ──────────────── 3s 11ms/step - loss: 0.0944 - val_loss: 0.0932
```

Test Loss: 0.09166844934225082
Test Accuracy: 0.9712756377551021


```python
# Display reconstruction with threshold
ax = plt.subplot(2, n, i + 1 + n)
reconstruction = decoded_imgs[i].reshape(28, 28)
plt.imshow(np.where(reconstruction >= threshold, 1.0, 0.0))
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)
plt.show()
```


Result

Thus, an Autoencoder has been implemented using Keras / Tensorflow.

Ex No:8   OBJECT DETECTION WITH YOLO3

DATE:29/09/2025

Aim:

To build an object detection model with YOLO3 using Keras/TensorFlow.

Procedure:

1. Download and load the dataset.

2. Perform analysis and preprocessing of the dataset.

3. Build a simple neural network model using Keras/TensorFlow.

4. Compile and fit the model.

5. Perform prediction with the test dataset.

6. Calculate performance metrics.

Code:
```
import cv2
import matplotlib.pyplot as plt
import numpy as np

# Define the paths to the YOLOv3 configuration, weights, and class names
files cfg_file = '/content/yolov3.cfg'
weight_file = '/content/yolov3.weights'
namesfile = '/content/coco.names'

# Load the YOLOv3 model
net = cv2.dnn.readNet(weight_file, cfg_file)

# Load class names
with open(namesfile, 'r') as f:
classes = f.read().strip().split('\n')

# Load an image for object
detection image_path =
'/content/hit.jpg' image =
cv2.imread(image_path)

# Get the height and width of the image
height, width = image.shape[:2]

# Create a blob from the image
blob = cv2.dnn.blobFromImage(image, 1/255.0, (416, 416), swapRB=True, crop=False)
net.setInput(blob)

# Get the names of the output layers
layer_names = net.getUnconnectedOutLayersNames()
```

Run forward pass

```python
outs = net.forward(layer_names)

# Initialize lists to store detected objects' information
class_ids = []
confidences = []
boxes = []

# Define a confidence threshold for object
detection conf_threshold = 0.5

# Loop over the detections
for out in outs:
for detection in out:
scores = detection[5:]
class_id =
np.argmax(scores)
confidence =
scores[class_id]
if confidence > conf_threshold:
# Object detected
center_x = int(detection[0] * width)
center_y  =  int(detection[1]  *
height)  w  =  int(detection[2]  *
width)
h = int(detection[3] * height)

#            Rectangle
coordinates   x   =
int(center_x - w / 2) y =
int(center_y - h / 2)

class_ids.append(class_id)
confidences.append(float(confidence)
) boxes.append([x, y, w, h])

# Apply non-maximum suppression to eliminate overlapping boxes
nms_threshold = 0.4
indices = cv2.dnn.NMSBoxes(boxes, confidences, conf_threshold, nms_threshold)

# Draw bounding boxes and labels on the image
for i in indices.flatten(): # flatten for
compatibility x, y, w, h = boxes[i]
label = str(classes[class_ids[i]])
confidence = confidences[i]
```

```
cv2.rectangle(image, (x, y), (x + w, y + h), (0, 255, 0), 2)
cv2.putText(image, f'{label} {confidence:.2f}', (x, y - 10),
cv2.FONT_HERSHEY_SIMPLEX, 0.8, (0, 255, 0), 2)
```

```
# Display the result in Jupyter Notebook
plt.figure(figsize=(10, 8))
plt.imshow(cv2.cvtColor(image, cv2.COLOR_BGR2RGB))
plt.axis('off')
plt.show()
```

Output:

```
cfg_file=r"C:\Users\agaki\Downloads\yolov3.cfg"
weight_file=r"C:\Users\agaki\Downloads\yolov3.weights"
names_file=r"C:\Users\agaki\Downloads\coco.names"
```

```
for out in outs:
    for detection in out:
        scores=detection[5:]
        class_id=np.argmax(scores)
        confidence=scores[class_id]
        if confidence>conf_threshold:

            center_x=int(detection[0]*width)
            center_y=int(detection[1]*height)
            w=int(detection[2]*width)
            h=int(detection[3]*height)


            x=int(center_x-w/2)
            y=int(center_y-h/2)

            class_ids.append(class_id)
            confidences.append(float(confidence))
            boxes.append([x,y,w,h])
```


Result

Thus, object detection using YOLOV5 has been implemented successfully.

Ex No: 9   BUILD GENERATIVE ADVERSARIAL NEURAL NETWORK

DATE:29/09/2025

Aim:

To build a generative adversarial neural network using Keras/TensorFlow.

Procedure:

1. Download and load the dataset.

2. Perform analysis and preprocessing of the dataset.

3. Build a simple neural network model using Keras/TensorFlow.

4. Compile and fit the model.

5. Perform prediction with the test dataset.

6. Calculate performance metrics.

Code:

```
import numpy as np
import tensorflow as tf
from tensorflow.keras.layers import Dense
from tensorflow.keras.models import Sequential
from tensorflow.keras.optimizers import Adam
from sklearn.datasets import load_iris
import matplotlib.pyplot as plt

# Load and Preprocess the Iris Dataset
iris = load_iris()
x_train = iris.data

# Build the GAN model
def build_generator():
model = Sequential()
model.add(Dense(128, input_shape=(100,), activation='relu'))
model.add(Dense(4, activation='linear')) # Output 4 features
return model

def build_discriminator():
model = Sequential()
model.add(Dense(128, input_shape=(4,), activation='relu'))
model.add(Dense(1, activation='sigmoid'))
return model

def build_gan(generator, discriminator):
discriminator.trainable = False
model = Sequential()
model.add(generator)
model.add(discriminator)
```

```python
    return model

generator = build_generator()
discriminator = build_discriminator()
gan = build_gan(generator, discriminator)

# Compile the Models
generator.compile(loss='mean_squared_error', optimizer=Adam(0.0002, 0.5))
discriminator.compile(loss='binary_crossentropy', optimizer=Adam(0.0002, 0.5),
metrics=['accuracy'])
gan.compile(loss='binary_crossentropy', optimizer=Adam(0.0002, 0.5))

# Training Loop
epochs = 200
batch_size = 16

for epoch in range(epochs):
# Train discriminator
idx = np.random.randint(0, x_train.shape[0], batch_size)
real_samples = x_train[idx]
fake_samples = generator.predict(np.random.normal(0, 1, (batch_size, 100)), verbose=0)

real_labels = np.ones((batch_size, 1))
fake_labels = np.zeros((batch_size, 1))

d_loss_real = discriminator.train_on_batch(real_samples, real_labels)
d_loss_fake = discriminator.train_on_batch(fake_samples, fake_labels)

# Train generator
noise = np.random.normal(0, 1, (batch_size, 100))
g_loss = gan.train_on_batch(noise, real_labels)

# Print progress
print(f"Epoch {epoch}/{epochs} | Discriminator Loss: {0.5 * (d_loss_real[0] + d_loss_fake[0])} |
Generator Loss: {g_loss}")

# Generating Synthetic Data
synthetic_data = generator.predict(np.random.normal(0, 1, (150, 100)), verbose=0)

# Create scatter plots for feature pairs
plt.figure(figsize=(12, 8))
plot_idx = 1

for i in range(4):
for j in range(i + 1, 4):
plt.subplot(2, 3, plot_idx)
```

```python
plt.scatter(x_train[:, i], x_train[:, j], label='Real Data', c='blue', marker='o', s=30)
plt.scatter(synthetic_data[:, i], synthetic_data[:, j], label='Synthetic Data', c='red', marker='x', s=30)
plt.xlabel(f'Feature {i + 1}')
plt.ylabel(f'Feature {j + 1}')
plt.legend()
plot_idx += 1

plt.tight_layout()
plt.show()
```

Output:

```
Epoch 0/200 | Discriminator Loss: 0.8773080408573151 |Generator Loss: 0.764731228351593
Epoch 1/200 | Discriminator Loss: 0.9332943856716156 |Generator Loss: 0.7988691329956055
Epoch 2/200 | Discriminator Loss: 0.9277275502681732 |Generator Loss: 0.8127573728561401
Epoch 3/200 | Discriminator Loss: 0.8921994566917419 |Generator Loss: 0.7757299542427063
Epoch 4/200 | Discriminator Loss: 0.913447916507721 |Generator Loss: 0.7737997174263
Epoch 5/200 | Discriminator Loss: 0.8916181325912476 |Generator Loss: 0.8003895282745361
Epoch 6/200 | Discriminator Loss: 0.9026078879833221 |Generator Loss: 0.814433217048645
Epoch 7/200 | Discriminator Loss: 0.9135120809078217 |Generator Loss: 0.8237183690071106
Epoch 8/200 | Discriminator Loss: 0.879832923412323 |Generator Loss: 0.7563657760620117
Epoch 9/200 | Discriminator Loss: 0.9439513385295868 |Generator Loss: 0.7623365521430969
Epoch 10/200 | Discriminator Loss: 0.9355685114860535 |Generator Loss: 0.7924684286117554
Epoch 11/200 | Discriminator Loss: 0.9386743903160095 |Generator Loss: 0.7614541053771973
Epoch 12/200 | Discriminator Loss: 0.960555225610733 |Generator Loss: 0.7792538404464722
Epoch 13/200 | Discriminator Loss: 0.9134297668933868 |Generator Loss: 0.792992115020752
Epoch 14/200 | Discriminator Loss: 0.8851655125617981 |Generator Loss: 0.7628173232078552
Epoch 15/200 | Discriminator Loss: 0.9505723416805267 |Generator Loss: 0.7851851582527161
Epoch 16/200 | Discriminator Loss: 0.92226842045784 |Generator Loss: 0.769191563129425
Epoch 17/200 | Discriminator Loss: 0.8982412815093994 |Generator Loss: 0.7685977220535278
Epoch 18/200 | Discriminator Loss: 0.9125983119010925 |Generator Loss: 0.7730982899665833
Epoch 19/200 | Discriminator Loss: 0.9367325305938721 |Generator Loss: 0.7837406396865845
Epoch 20/200 | Discriminator Loss: 0.9531015455722809 |Generator Loss: 0.7827053070068359
Epoch 21/200 | Discriminator Loss: 0.9306998252868652 |Generator Loss: 0.7667914032936096
Epoch 22/200 | Discriminator Loss: 0.8887360095977783 |Generator Loss: 0.7845874428749084
Epoch 23/200 | Discriminator Loss: 0.9426513016223907 |Generator Loss: 0.746765673160553
Epoch 24/200 | Discriminator Loss: 0.9331325888633728 |Generator Loss: 0.761589765548706
Epoch 25/200 | Discriminator Loss: 0.9080778360366821 |Generator Loss: 0.7709233164787292
Epoch 26/200 | Discriminator Loss: 0.9232879281044006 |Generator Loss: 0.7773635387420654
Epoch 27/200 | Discriminator Loss: 0.9102294743061066 |Generator Loss: 0.7809370756149292
Epoch 28/200 | Discriminator Loss: 0.9312145709991455 |Generator Loss: 0.7647197246551514
Epoch 29/200 | Discriminator Loss: 0.9415165781974792 |Generator Loss: 0.7561923861503601
Epoch 30/200 | Discriminator Loss: 0.930676281452179 |Generator Loss: 0.7709008455276489
Epoch 31/200 | Discriminator Loss: 0.9495892226696014 |Generator Loss: 0.7595088481903076
```



```python
synthetic_data = generator.predict(np.random.normal(0,1,(150,100)),verbose=0)
plt.figure(figsize=(12,8))
plot_idx=1

for i in range(4):
    for j in range(i+1,4):
        plt.subplot(2,3,plot_idx)
        plt.scatter(x_train[:,i],x_train[:,j],label='Real Data',c='blue',marker='o',s=30)
        plt.scatter(synthetic_data[:,i],synthetic_data[:,j],label='Synthetic Data',c='red',marker='x',s=30)
        plt.xlabel(f'Feature{i+1}')
        plt.ylabel(f'Feature{j+1}')
        plt.legend()
        plot_idx+=1

plt.tight_layout()
plt.show()
```
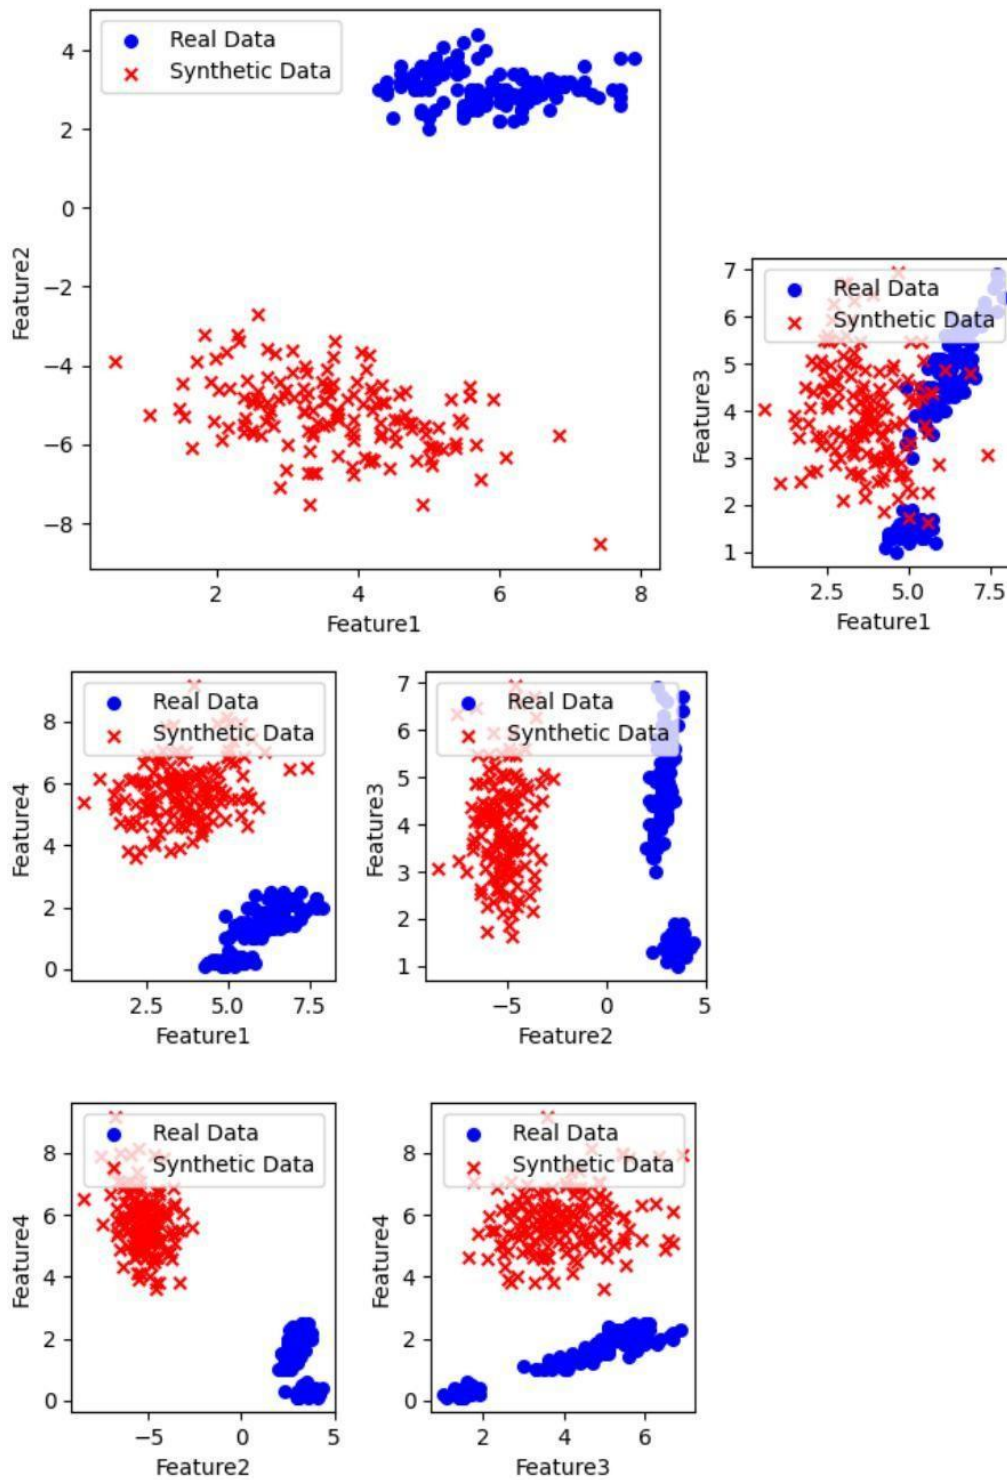
Result

Thus, a generative adversarial neural network using Keras / Tensorflow has been implemented successfully.

Ex No:10

DATE:06/10/2025

MINI PROJECT

BreedVision

Aim:

Dog Breed Identification using Deep Learning (ResNet18).

Code:

```python
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader
from torchvision import datasets, transforms, models

import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import classification_report, confusion_matrix
from tqdm import tqdm
import os
import warnings
warnings.filterwarnings('ignore')

# Check GPU
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f"PyTorch Version: {torch.__version__}")
print(f"Device: {device}")
if torch.cuda.is_available():
    print(f"GPU: {torch.cuda.get_device_name(0)}")
    print(f"GPU Memory: {torch.cuda.get_device_properties(0).total_memory / 1024**3:.2f} GB")
```

```python
# Paths
DATA_DIR = 'dataset_small'  # Change to 'datase

# Hyperparameters
BATCH_SIZE = 64
NUM_EPOCHS = 50
LEARNING_RATE = 0.001
IMG_SIZE = 224  # ResNet expects 224x224

# Set random seeds for reproducibility
torch.manual_seed(42)
np.random.seed(42)

print(f"Dataset: {DATA_DIR}")
print(f"Batch Size: {BATCH_SIZE}")
print(f"Epochs: {NUM_EPOCHS}")
print(f"Learning Rate: {LEARNING_RATE}")
print(f"Image Size: {IMG_SIZE}x{IMG_SIZE}")
```

```python
# Load pre-trained ResNet18
model = models.resnet18(pretrained=True)

# Freeze early layers
for param in model.parameters():
    param.requires_grad = False

# Replace final layer
num_features = model.fc.in_features
model.fc = nn.Sequential(
    nn.Linear(num_features, 512),
    nn.ReLU(),
    nn.Dropout(0.5),
    nn.Linear(512, 256),
    nn.ReLU(),
    nn.Dropout(0.3),
    nn.Linear(256, num_classes)
)

model = model.to(device)

# Loss and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.fc.parameters(), lr=LEARNING_RATE)
scheduler = optim.lr_scheduler.ReduceLROnPlateau(optimizer, mode='min', patience=3, factor=0.5)

print("Model Architecture:")
print(f"Using ResNet18 with transfer learning")
print(f"Trainable parameters: {sum(p.numel() for p in model.parameters() if p.requires_grad):,}")
print(f"Total parameters: {sum(p.numel() for p in model.parameters()):,}")
```

```python
# Training Loop
history = {
    'train_loss': [],
    'train_acc': [],
    'val_loss': [],
    'val_acc': []
}

best_val_acc = 0.0

print("="*60)
print("STARTING TRAINING")
print("="*60)

for epoch in range(NUM_EPOCHS):
    print(f"\nEpoch {epoch+1}/{NUM_EPOCHS}")
    print("-" * 60)

    # Train
    train_loss, train_acc = train_epoch(model, train_loader, criterion, optimizer, device)

    # Validate
    val_loss, val_acc = validate(model, val_loader, criterion, device)

    # Update learning rate
    scheduler.step(val_loss)

    # Save history
    history['train_loss'].append(train_loss)
    history['train_acc'].append(train_acc)
    history['val_loss'].append(val_loss)
    history['val_acc'].append(val_acc)

    print(f"\nTrain Loss: {train_loss:.4f} | Train Acc: {train_acc:.2f}%")
    print(f"Val Loss: {val_loss:.4f} | Val Acc: {val_acc:.2f}%")

    # Save best model
    if val_acc > best_val_acc:
        best_val_acc = val_acc
        torch.save(model.state_dict(), 'best_model.pth')
        print(f"√ Best model saved! (Val Acc: {val_acc:.2f}%)")

print("\n" + "="*60)
print("TRAINING COMPLETED!")
print("="*60)
print(f"Best Validation Accuracy: {best_val_acc:.2f}%")
```

```python
# Get a batch of test images
images, labels = next(iter(test_loader))
images = images.to(device)

model.eval()
with torch.no_grad():
    outputs = model(images)
    probs = torch.nn.functional.softmax(outputs, dim=1)
    _, preds = outputs.max(1)

# Plot
fig, axes = plt.subplots(3, 3, figsize=(15, 15))
axes = axes.ravel()

for idx in range(9):
    img = imshow(images[idx].cpu())
    axes[idx].imshow(img)

    true_label = class_names[labels[idx]]
    pred_label = class_names[preds[idx]]
    confidence = probs[idx][preds[idx]].item() * 100

    color = 'green' if labels[idx] == preds[idx] else 'red'
    axes[idx].set_title(f"True: {true_label}\nPred: {pred_label}\nConf: {confidence:.1f}%",
                        color=color, fontsize=11, fontweight='bold')
    axes[idx].axis('off')

plt.suptitle('Model Predictions (Green=Correct, Red=Incorrect)',
             fontsize=16, fontweight='bold')
plt.tight_layout()
plt.show()
```

```python
## 13. Predict Custom Dog Images

### Test on Any Dog Image!

def predict_custom_image(image_path, model, class_names, device):
    """
    Predict dog breed from any custom image

    Args:
        image_path: Path to your dog image (jpg, png, etc.)
        model: Trained PyTorch model
        class_names: List of breed names
        device: cuda or cpu
    """
    from PIL import Image

    # Check if file exists
    if not os.path.exists(image_path):
        print(f"Error: Image not found at {image_path}")
        return

    # Load and display original image
    image = Image.open(image_path).convert('RGB')

    # Transform for model
    transform = transforms.Compose([
        transforms.Resize((224, 224)),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
    ])

    image_tensor = transform(image).unsqueeze(0).to(device)

    # Predict
    model.eval()
    with torch.no_grad():
        outputs = model(image_tensor)
        probs = torch.nn.functional.softmax(outputs, dim=1)
        top5_probs, top5_indices = torch.topk(probs, min(5, len(class_names)))

    # Get results
    top5_probs = top5_probs[0].cpu().numpy()
    top5_indices = top5_indices[0].cpu().numpy()

    # Display results
    fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 6))

    # Show original image
    ax1.imshow(image)
    ax1.axis('off')
    predicted_breed = class_names[top5_indices[0]]
```

```python
    predicted_breed = class_names[top5_indices[0]]
    confidence = top5_probs[0] * 100
    ax1.set_title(f'Predicted Breed: {predicted_breed}\nConfidence: {confidence:.1f}%',
                  fontsize=14, fontweight='bold', color='green')

    # Show top 5 predictions as bar chart
    breeds = [class_names[idx] for idx in top5_indices]
    confidences = top5_probs * 100

    colors = ['green', 'orange', 'yellow', 'lightblue', 'lightcoral']
    bars = ax2.barh(breeds, confidences, color=colors[:len(breeds)])
    ax2.set_xlabel('Confidence (%)', fontsize=12)
    ax2.set_title('Top 5 Breed Predictions', fontsize=14, fontweight='bold')
    ax2.set_xlim(0, 100)

    # Add percentage labels
    for i, (breed, conf) in enumerate(zip(breeds, confidences)):
        ax2.text(conf + 2, i, f'{conf:.1f}%', va='center', fontsize=10, fontweight='bold')

    plt.tight_layout()
    plt.show()

    # Print detailed results
    print("\n" + "="*70)
    print("PREDICTION RESULTS")
    print("="*70)
    print(f"\n🏆 Top Prediction: {predicted_breed}")
    print(f"📊 Confidence: {confidence:.2f}%")
    print(f"\n📋 Top 5 Predictions:")
    for i, (idx, prob) in enumerate(zip(top5_indices, top5_probs), 1):
        print(f"  {i}. {class_names[idx]:<30} {prob*100:>6.2f}%")
    print("="*70)

    return predicted_breed, confidence

# Load best model
print("Loading best model...")
model.load_state_dict(torch.load('best_model.pth'))
model.eval()
print("√ Model loaded successfully!")
```

Test Single Image

```
# Replace 'path/to/your/dog.jpg' with your actual image path
image_path = 'dog.jpeg'
predict_custom_image(image_path, model, class_names, device)
```

Test Multiple Images

```python
import glob

# Put all your test images in a folder
test_folder = 'my_test_images'  # Create this folder and add images

# Get all images
image_paths = glob.glob(f'{test_folder}/*.jpg') + glob.glob(f'{test_folder}/*.png')

print(f"Found {len(image_paths)} images to test\n")

for img_path in image_paths:
    print(f"\n{'='*70}")
    print(f"Testing: {img_path}")
    print('='*70)
    predict_custom_image(img_path, model, class_names, device)
```

Test Random from Test Set

```python
import random

# Get a random test image
test_samples = test_dataset.samples
random_sample = random.choice(test_samples)
image_path, true_label = random_sample

print(f"True breed: {class_names[true_label]}")
predict_custom_image(image_path, model, class_names, device)
```

```python
# Save complete model
torch.save({
    'model_state_dict': model.state_dict(),
    'optimizer_state_dict': optimizer.state_dict(),
    'class_names': class_names,
    'num_classes': num_classes,
    'best_val_acc': best_val_acc,
    'test_acc': test_acc
}, 'dog_breed_model_complete.pth')

print("✓ Model saved as 'dog_breed_model_complete.pth'")
print("✓ Best model saved as 'best_model.pth'")
```

Output:

```
============================================================
STARTING TRAINING
============================================================

Epoch 1/50
------------------------------------------------------------
Training: 100%|████████████████████████████| 25/25 [00:15<00:00,  1.61it/s, loss=1.57, acc=45.7]
Validation: 100%|██████████████████████████| 6/6 [00:07<00:00,  1.17s/it]

Train Loss: 1.5730 | Train Acc: 45.73%
Val Loss: 0.4909 | Val Acc: 88.72%
√ Best model saved! (Val Acc: 88.72%)

Epoch 2/50
------------------------------------------------------------
Training: 100%|████████████████████████████| 25/25 [00:15<00:00,  1.64it/s, loss=0.674, acc=75.4]
Validation: 100%|██████████████████████████| 6/6 [00:07<00:00,  1.17s/it]

Train Loss: 0.6739 | Train Acc: 75.37%
Val Loss: 0.4042 | Val Acc: 84.45%

Epoch 3/50
------------------------------------------------------------
Training: 100%|████████████████████████████| 25/25 [00:15<00:00,  1.61it/s, loss=0.537, acc=80.8]
Validation: 100%|██████████████████████████| 6/6 [00:07<00:00,  1.21s/it]

Train Loss: 0.5367 | Train Acc: 80.82%
Val Loss: 0.2880 | Val Acc: 88.11%

Epoch 4/50
------------------------------------------------------------
Training: 100%|████████████████████████████| 25/25 [00:16<00:00,  1.54it/s, loss=0.392, acc=85.2]
Validation: 100%|██████████████████████████| 6/6 [00:07<00:00,  1.20s/it]

Train Loss: 0.3917 | Train Acc: 85.25%
Val Loss: 0.2597 | Val Acc: 87.80%

Epoch 5/50
------------------------------------------------------------
Training: 100%|████████████████████████████| 25/25 [00:15<00:00,  1.61it/s, loss=0.38, acc=87.1]
Validation: 100%|██████████████████████████| 6/6 [00:07<00:00,  1.23s/it]

Train Loss: 0.3796 | Train Acc: 87.11%
Val Loss: 0.2670 | Val Acc: 90.55%
√ Best model saved! (Val Acc: 90.55%)

Epoch 6/50
------------------------------------------------------------
Training: 100%|████████████████████████████| 25/25 [00:14<00:00,  1.67it/s, loss=0.38, acc=86.1]
Validation: 100%|██████████████████████████| 6/6 [00:06<00:00,  1.14s/it]
```
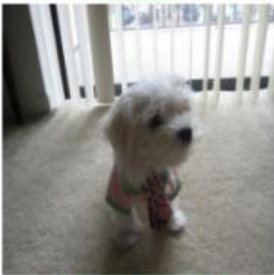


```
Validation: 100%|████████████████████████████| 6/6 [00:09<00:00,  1.56s/it]

============================================================
TEST SET RESULTS
============================================================
Test Loss: 0.2549
Test Accuracy: 90.75%
============================================================
```

Confusion Matrix

Model Predictions (Green=Correct, Red=Incorrect)

Predicted: n02088094-Afghan_hound
Confidence: 100.0%

Top 5 Predictions

Predicted: n02111889-Samoyed
Confidence: 99.3%

Top 5 Predictions

Predicted: n02112018-Pomeranian
Confidence: 99.9%

Top 5 Predictions

Result:

Thus, an Deep Learning model to identify different dog breeds has been implemented successfully.